

F# for the Analyst and the High-Performance Engineer

By Matthew Crews

SDI

Who am I?



F# for the Analyst

Defining Types is Easy

```
type Chicken =  
{  
    Name: string  
    Size: float  
    Age: int  
}
```

```
type Animal =  
    | Chicken of Chicken  
    | Turkey of Turkey  
    | Goose of Goose
```

Type Inference

```
let myFunction a b c = // int -> int -> int  
|...| a + b / c
```

```
module Chicken =  
  
  let create name age size = // string -> int -> float -> Chicken  
  |...| {  
    |...|   Name = name  
    |...|   Age = age  
    |...|   Size = size  
    |...| }
```

Pattern Matching

```
let getName a = // Animal -> string
  match a with
  | Chicken Incomplete pattern matches on this expression. F# Compiler(25)
  | Turkey
  | Goose g
    Open the documentation
    Full name: x
let checkBound Assembly: CodeExamples
  match x w
  | a when a < 0.0 -> negative
    View Problem (Alt+F8) No quick fixes available
```

Type Providers

```
1  Name, Age, Size
2  Clucky, 1, 2.0
3  Gobble, 10, 14.0
4  Scramble, 2, 3.5

#r "nuget: FSharp.Data"
open FSharp.Data

type ChickenRow = CsvProvider<"ExampleData.csv">

let chickens = // CsvProvider<...>
|> ChickenRow.Load("ExampleData.csv")

for row in chickens.Rows do
|> printfn $"{row.Name}"
```

Computation Expressions

```
let model =  
  Model "Coffee" {  
    structure [  
      supply1 → cartoner1  
      supply2 → cartoner2  
      supply3 → cartoner3  
      supply4 → cartoner4  
      cartoner1 → cartonerMerge  
      cartoner2 → cartonerMerge  
      cartoner3 → cartonerMerge  
      cartoner4 → cartonerMerge  
      cartonerMerge → casePackerFeed  
      casePackerFeed → casePacker  
      casePacker → palletizerFeed  
      palletizerFeed → palletizer  
      palletizer → wrapperFeed  
      wrapperFeed → wrapper  
      wrapper → endingInventory  
    ]  
    limits [  
      cartoner1, 2.0  
      cartoner2, 2.0  
      cartoner3, 2.0  
      cartoner4, 2.0  
      casePacker, 2.3  
      palletizer, 2.6  
      wrapper, 2.9  
    ]  
    capacities [  
      supply1, infVolume  
      supply2, infVolume  
      supply3, infVolume  
      supply4, infVolume  
      casePackerFeed, 10.0  
      palletizerFeed, 10.0  
      wrapperFeed, 10.0  
      endingInventory, infVolume  
    ]  
    levels [  
      supply1, infVolume  
      supply2, infVolume  
      supply3, infVolume  
      supply4, infVolume  
    ]  
    mergeRules [  
      cartonerMerge, (MergeRule.Priority [  
        cartoner1; cartoner2; cartoner3; cartoner4  
      ])  
    ]
```

Create .NET apps faster with NuGet

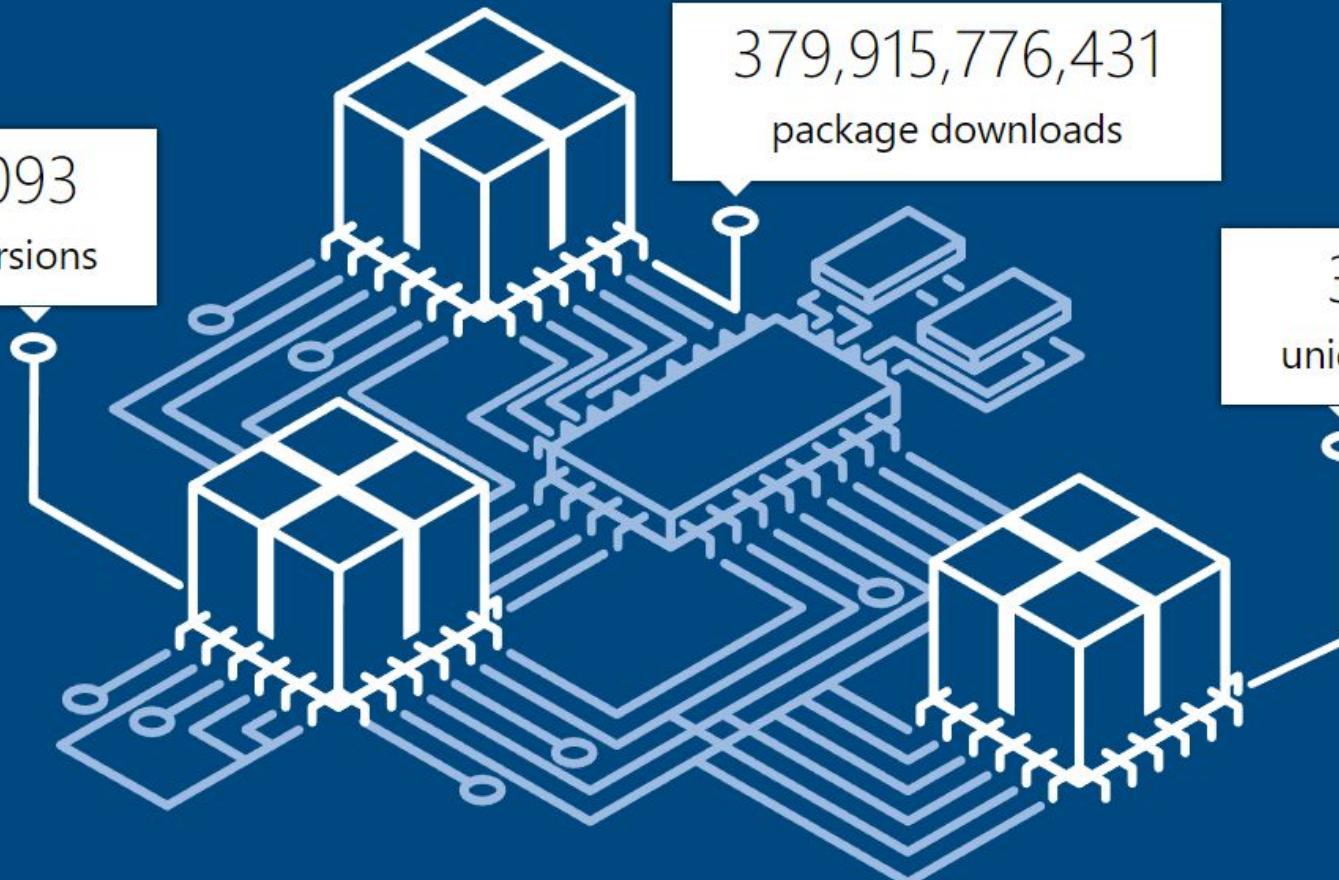
Search for packages...



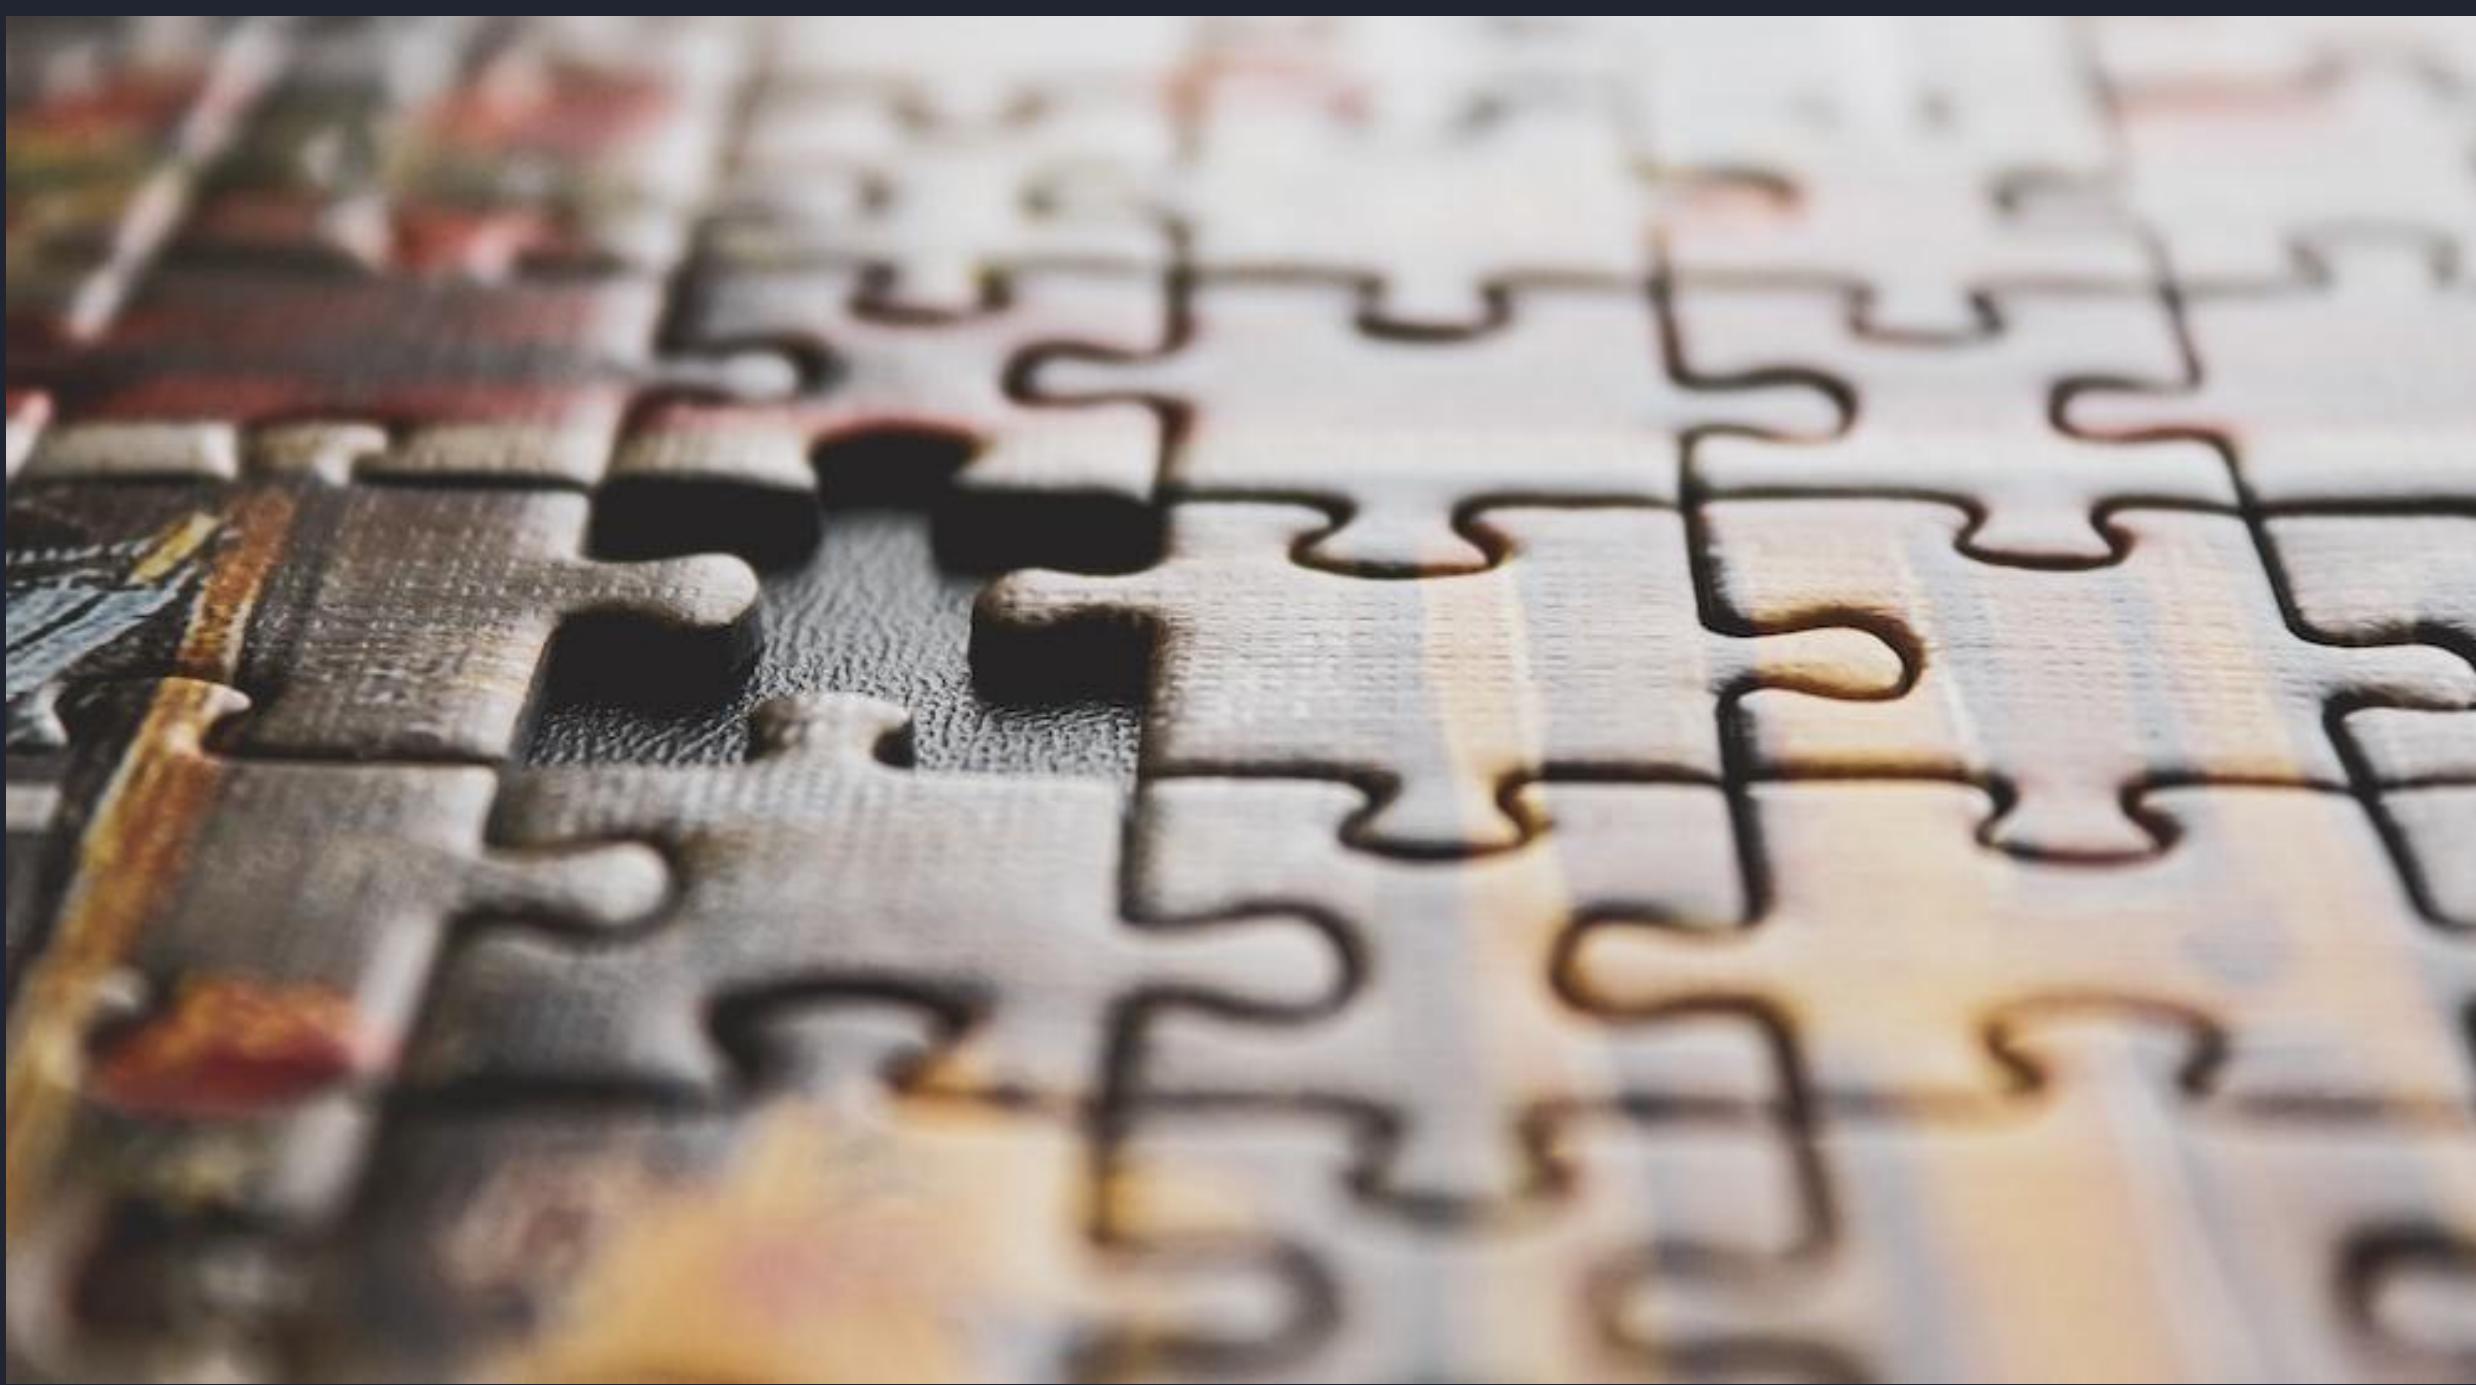
6,254,093
package versions

379,915,776,431
package downloads

371,706
unique packages









Innovation Rate

Time to Develop x Time to Execute

F# for High-Performance

All of performance comes down to...

Doing Less, With Less, Predictably

Doing Less

Minimize the number of CPU instructions

F# is Statically Typed

F#

```
let add (a, b) =  
    a + b
```

ASM

```
_add(Int32, Int32)  
L0000: lea eax, [ecx+edx]  
L0003: ret
```

Python

```
def add(a, b)  
    return a + b
```

ASM

```
mov     qword ptr [r11+38h],r14  
add     r14,2  
jmp     _PyEval_EvalFrameDefault+5872h (07FFDFD04F462h)  
mov     rdi,qword ptr [r12-8]  
lea     rax,[__ImageBase (07FFDFCCE20000h)]  
mov     rsi,qword ptr [r12-10h]  
sub     r12,8  
mov     rdx,rdi  
mov     rcx,rsi  
call    qword ptr [rax+r13*8+3E7A70h]
```

```
mov    qword ptr [rsp+8],rbx
push   rdi
sub   rsp,30h
xor   r8d,r8d
mov   rdi,rdx
mov   rbx,rcx
call  binary_op1 (07FFDFCF06C70h)
mov   qword ptr [rsp+10h],rbp
mov   qword ptr [rsp+18h],rsi
mov   qword ptr [rsp+20h],rdi
push   r14
sub   rsp,20h
mov   rsi,rdx
movsxd r9,r8d
mov   rdx,qword ptr [rcx+8]
mov   rbp,rcx
mov   rdi,qword ptr [rdx+60h]
test  rdi,rdi
je   binary_op1+2Fh (07FFDFCF06C9Fh)
mov   rdi,qword ptr [rdi+r9]
mov   rcx,qword ptr [rsi+8]
mov   qword ptr [rsp+30h],rbx
cmp   rcx,rdx
je   binary_op1+55h (07FFDFCF06CC5h)
xor   ebx,ebx
lea   r14,[_Py_NotImplementedStruct (07FFDFD2E6AF0h)]
test  rdi,rdi
je   binary_op1+0ADh (07FFDFCF06D1Dh)
test  rbx,rbx
je   binary_op1+90h (07FFDFCF06D00h)
mov   rdx,rsi
mov   rcx,rbp
call  rdi
mov   rax,qword ptr [rcx+8]
test  dword ptr [rax+0A8h],1000000h
je   long_add+24h (07FFDFCF60434h)
mov   rax,qword ptr [rdx+8]

test  dword ptr [rax+0A8h],1000000h
jne   _PyLong_Add (07FFDFCF60380h)
sub   rsp,28h
mov   r8,rdx
mov   rdx,qword ptr [rcx+10h]
lea   rax,[rdx+1]
cmp   rax,3
jae  _PyLong_Add+3Eh (07FFDFCF603BEh)
mov   r9,qword ptr [r8+10h]
lea   rax,[r9+1]
cmp   rax,3
jae  _PyLong_Add+3Eh (07FFDFCF603BEh)
mov   ecx,dword ptr [rcx+18h]
mov   eax,dword ptr [r8+18h]
imul  rcx,rdx
imul  rax,r9
add   rcx,rcx
add   rsp,28h
jmp   _PyLong_FromSTwoDigits (07FFDFCF5AF10h)
push  rdi
sub   rsp,20h
lea   rax,[rcx+5]
mov   rdi,rcx
cmp   rax,105h
ja   _PyLong_FromSTwoDigits+2Fh (07FFDFCF5AF3Fh)
lea   rax,[rcx+3FFFFFFFh]
cmp   rax,7FFFFFFFh
jae  _PyLong_FromSTwoDigits+48h (07FFDFCF5AF58h)
add   rsp,20h
pop   rdi
jmp   _PyLong_FromMedium (07FFDFCF5AE50h)
mov   qword ptr [rsp+10h],rbx
push  rsi
sub   rsp,20h
movsxd rsi,ecx
mov   edx,20h
mov   rcx,qword ptr [_PyObject (07FFDFD3B8CA8h)]
```

Can you make Dynamic Typing fast?



PYTHON

Unladen Swallow [edit]

Unladen Swallow was an optimization branch of CPython
supplementing CPython's custom [virtual machine](#) with a

Mojo 🔥 – the
programming language

for all AI developers .



```
def sort(v: ArraySlice[Int]):  
    for i in range(len(v)):  
        for j in range(len(v) - i - 1):  
            if v[j] > v[j + 1]:  
                swap(v[j], v[j + 1])
```



```
struct MyPair:  
    var first: Int  
    var second: F32  
  
def __init__(self, first: Int, second: F32):  
    self.first = first  
    self.second = second
```

Default Equality

```
type Chicken =  
{  
    Name: string  
    Size: float  
    Age: int  
}  
  
let equalsTest (a: Chicken, b: Chicken) =  
    a = b
```

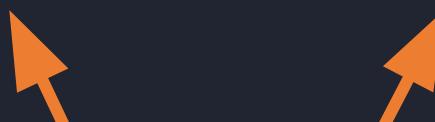
```
_.equalsTest(Chicken, Chicken)  
L0000: push ebp  
L0001: mov ebp, esp  
L0003: cmp [ecx], cl  
L0005: test edx, edx  
L0007: je short L003d  
L0009: mov ecx, [ecx+4]  
L000c: mov edx, [edx+4]  
L000f: cmp ecx, edx  
L0011: je short L0041  
L0013: test ecx, ecx  
L0015: je short L0023  
L0017: test edx, edx  
L0019: je short L0023  
L001b: mov eax, [ecx+4]  
L001e: cmp eax, [edx+4]  
L0021: je short L0027  
L0023: xor eax, eax  
L0025: jmp short L003f  
L0027: lea eax, [ecx+8]  
L002a: mov ecx, [ecx+4]  
L002d: add ecx, ecx  
L002f: push ecx  
L0030: add edx, 8  
L0033: mov ecx, eax  
L0035: call dword ptr [0x652d738]  
L003b: jmp short L0025  
L003d: xor eax, eax  
L003f: pop ebp  
L0040: ret  
L0041: mov eax, 1  
L0046: jmp short L0025
```

Units of Measure

```
[<Measure>]  
type ChickenId  
  
type ChickenData =  
| {  
|   Name: string  
|   Age: int  
|   Size: float  
| }
```

```
let equalsTestUoM (a: int<ChickenId>, b: int<ChickenId>) =  
| a = b
```

```
_.equalsTestUoM(Int32, Int32)  
L0000: xor eax, eax  
L0002: cmp ecx, edx  
L0004: sete al  
L0007: ret
```



We use numeric primitives to identify Entities.

Bounds Checking



C Devs

Bounds Checking

```
let test(a: int[]) = // []<int> -> int
  let mutable i = 0
  let mutable acc = 0
  while i < a.Length do
    acc <- acc + a[i]
  acc
```

```
_test(Int32[])
L0000: push ebp
L0001: mov ebp, esp
L0003: xor eax, eax
L0005: mov edx, [ecx+4]
L0008: test edx, edx
L000a: jle short L0011
L000c: add eax, [ecx+8]
L000f: jmp short L0008
L0011: pop ebp
L0012: ret
```



With Less

Minimize the memory footprint of your application

CPU Memory Architecture

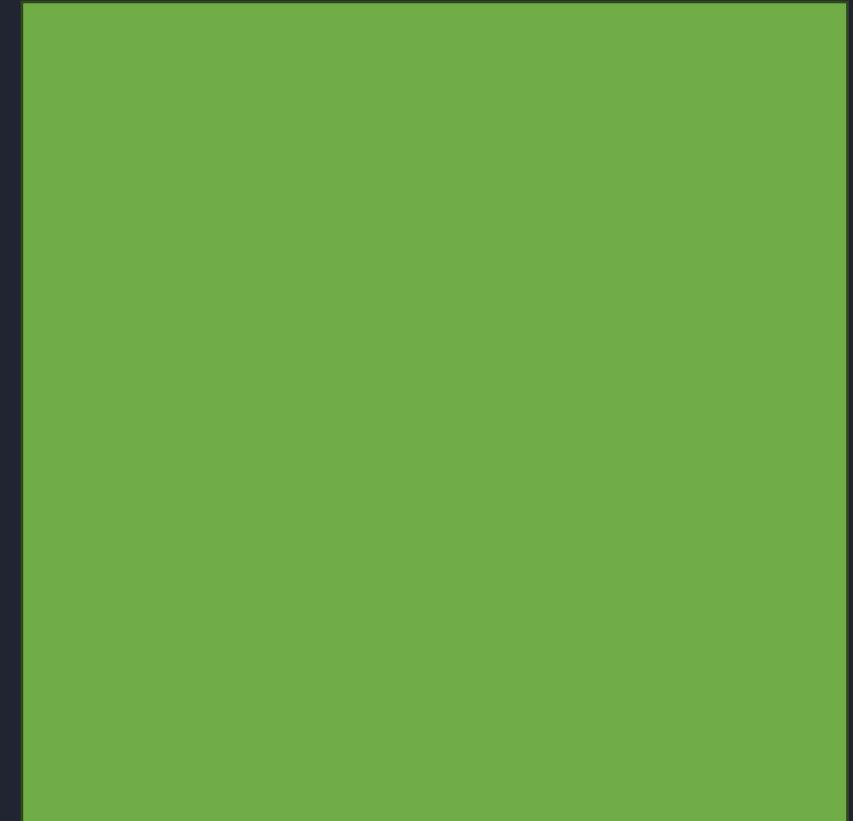
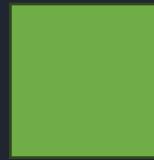
Registers
10-100
64 bit

L1 Cache
80 KB

L2 Cache
1.25 MB

L3 Cache
30 MB Shared

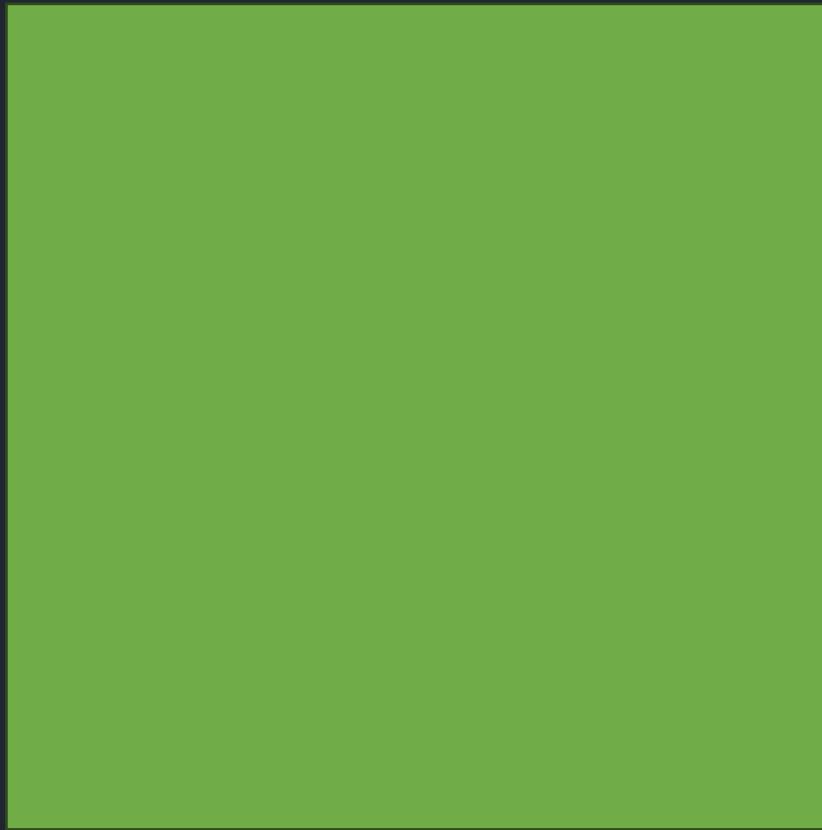
ALU/
FPU



CPU Memory Architecture

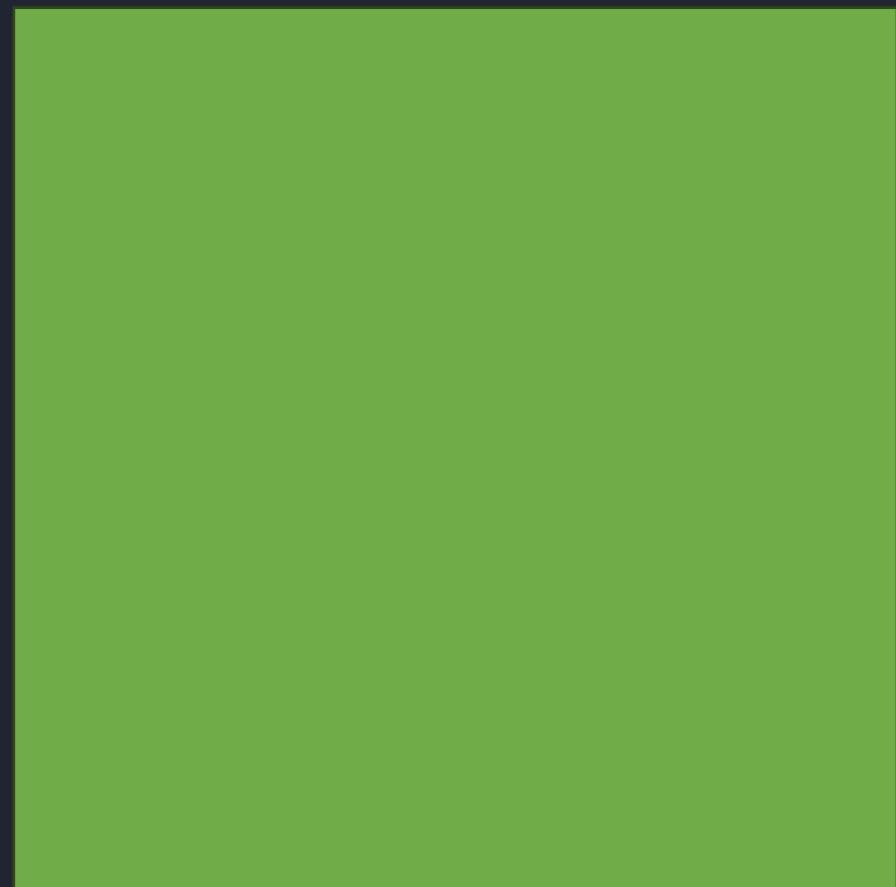
L3 Cache

30 MB



Main Memory

64 GB



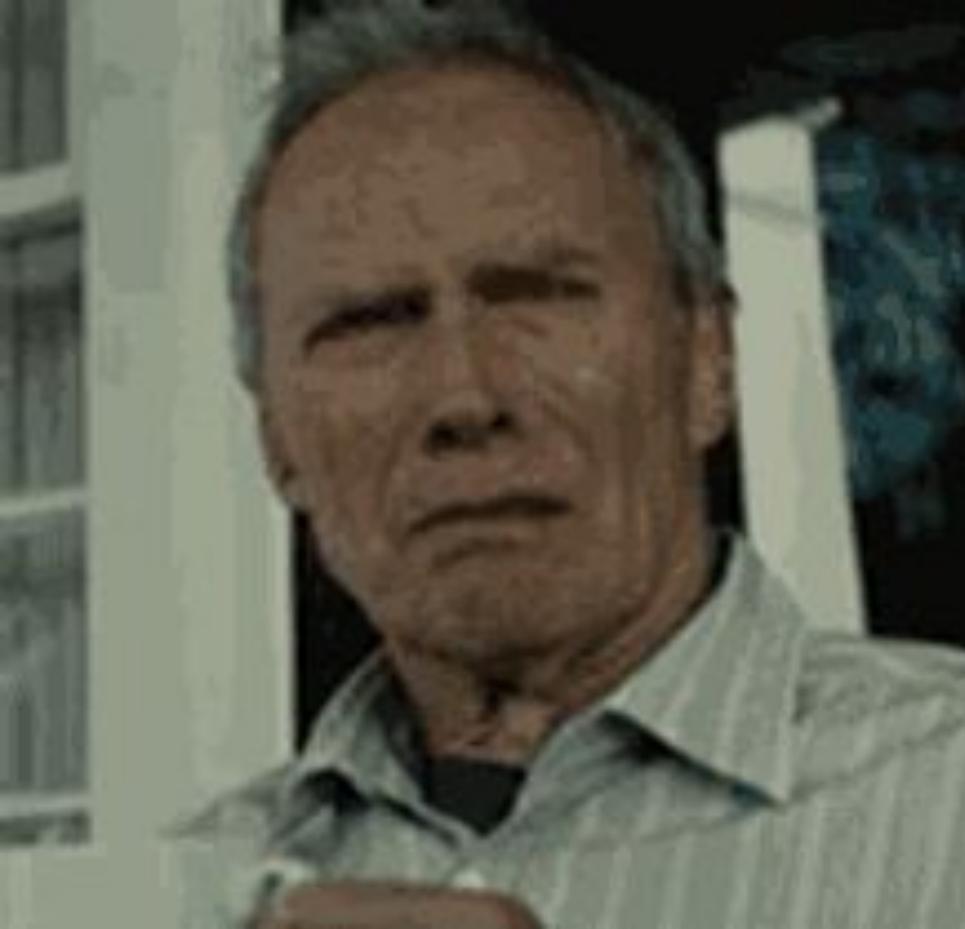
Memory Latency





Cache Space is incredibly
scarce

C Dev hearing your language has a
Garbage Collector...



What the GC Does

- Makes Allocating/Deallocating memory transparent
- Designed to be general-purpose and robust
- Makes writing horrendous code (for performance) easy

What the GC Does NOT do

- Understand your memory access patterns
- Does not co-locate data that is frequently accessed together
- Understand Allocation/Deallocation patterns (Arena Allocator)



Row

```
[<Struct>]
type Row<[<Measure>] 'Measure, 'T>(values: 'T[]) =
    [<>EditorBrowsable(EditorBrowsableState.Never)>]
    member __.values : 'T[] = values // []<'T>

    member r.Item
        with inline get (i: int<'Measure>) =
            r._values[int i]

        and inline set (index: int<'Measure>) value =
            r._values[int index] <- value

    member inline r.Length = LanguagePrimitives.Int32WithMeasure<'Measure> r._values.Length
```

Bar

```
[<Struct>]
type Bar<[<Measure>] 'Measure, 'T> (values: 'T[]) =
    [<>
        [<EditorBrowsable(EditorBrowsableState.Never)>]
        member __.values = values // []<'T>
    </>
    member inline b.Length = LanguagePrimitives.Int32WithMeasure<'Measure> b._values.Length
    member b.Item // 'T
        with inline get(i: int<'Measure>) =
            b._values[int i]
```

Row and Bar Usage

```
let · chickenAges · = · Bar<ChickenId, · _> · [ | 0; · 1; · 5; · 0; · 2 | ]  
let · chickenIndex · = · 1<ChickenId>  
// · The · Compiler · enforces · the · correct · units · on · the · indexing · Int  
let · chickenAge · = · chickenAges[chickenIndex]
```

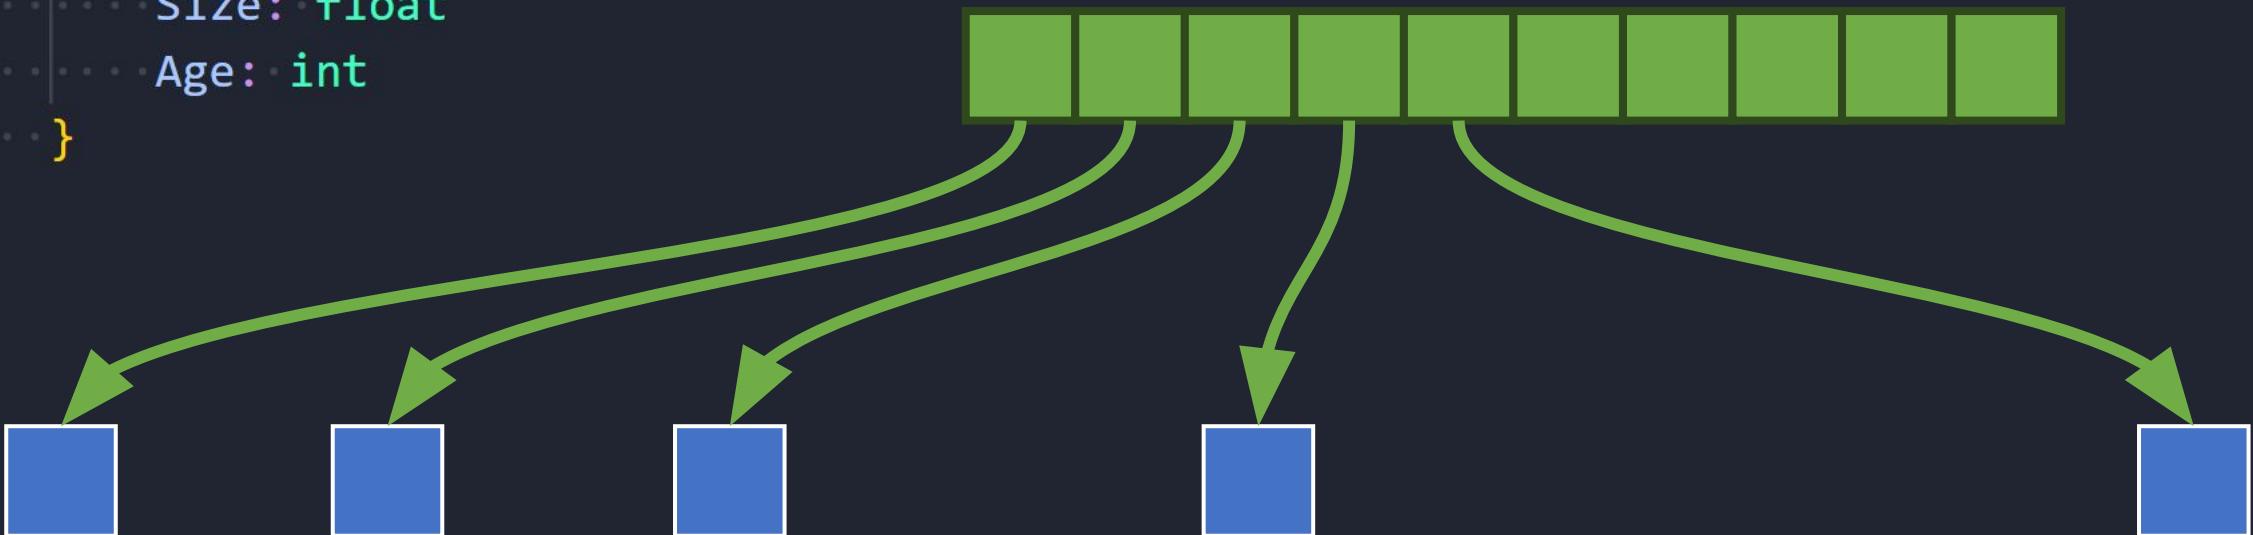
Row and Bar Usage

```
let · chickenAges · = · Bar<ChickenId, · _> · [ | 0; · 1; · 5; · 0; · 2 | ]  
let · chickenIndex · = · 1<ChickenId>  
// · The · Compiler · enforces · the · correct · units · on · the · indexing · Int  
let · chickenAge · = · chickenAges[chickenIndex]  
  
// · Non-Chicken · Index  
let · cowIndex · = · 1<CowId>  
// · This · a · compiler · error · since · the · units · of · the · indexing · value  
// · does · not · match · the · expectation · of · the · collection  
let · cowAge · = · chickenAges[cowIndex]
```

Array of Objects

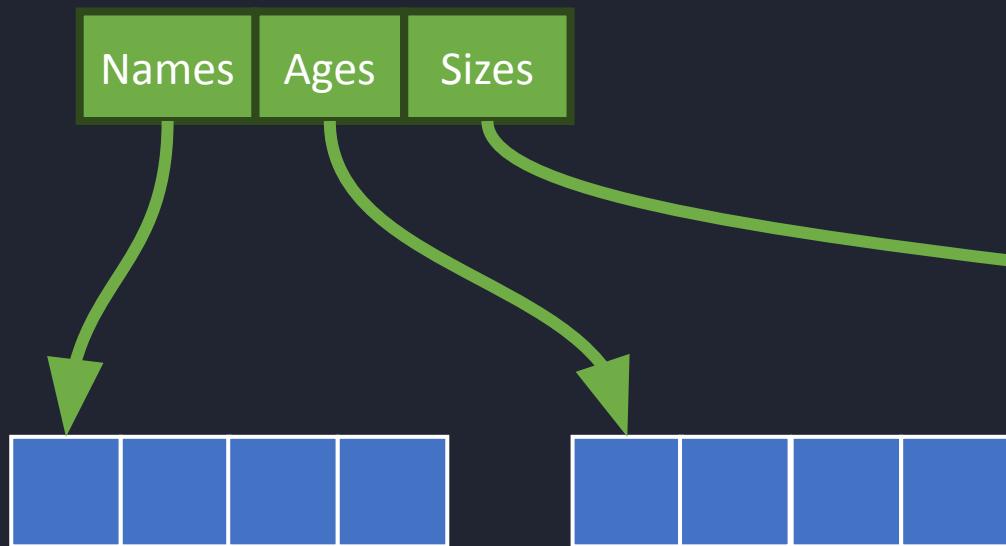
```
type Chicken =  
{  
    Name: string  
    Size: float  
    Age: int  
}
```

Chicken[]



Each element could be on a different Cache Line
and therefore incur a cache miss on load

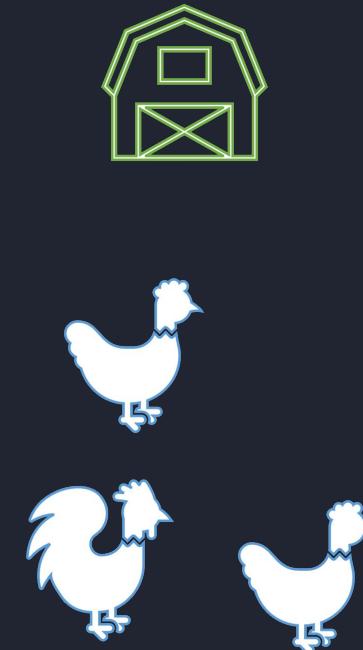
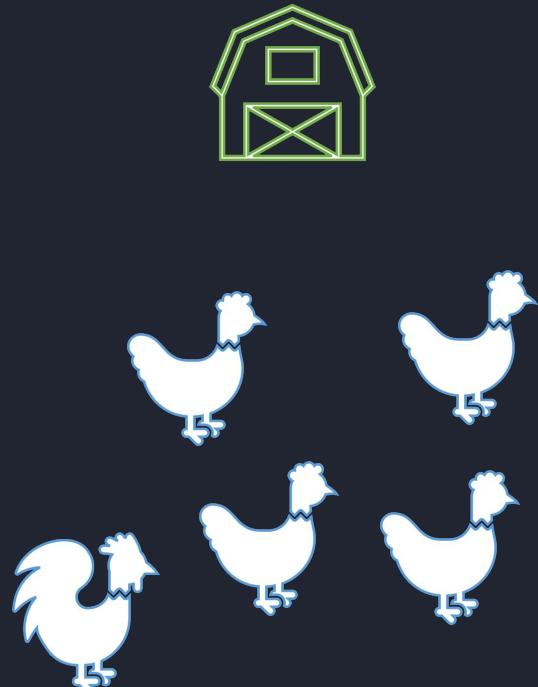
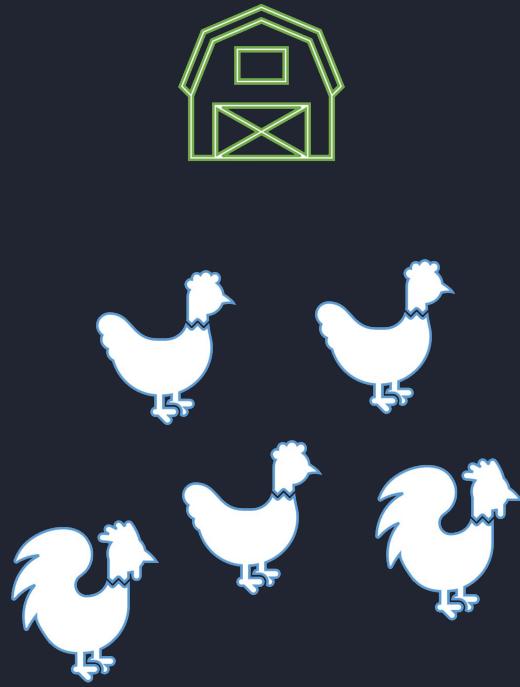
Struct of Arrays (SoA)



```
type Flock =  
{  
    ...  
    Names: Bar<ChickenId, string>  
    Ages: Row<ChickenId, int>  
    Sizes: Row<ChickenId, float>  
}
```

Values are contiguous in memory which minimizes the probability of a cache miss

One to Many Relationships



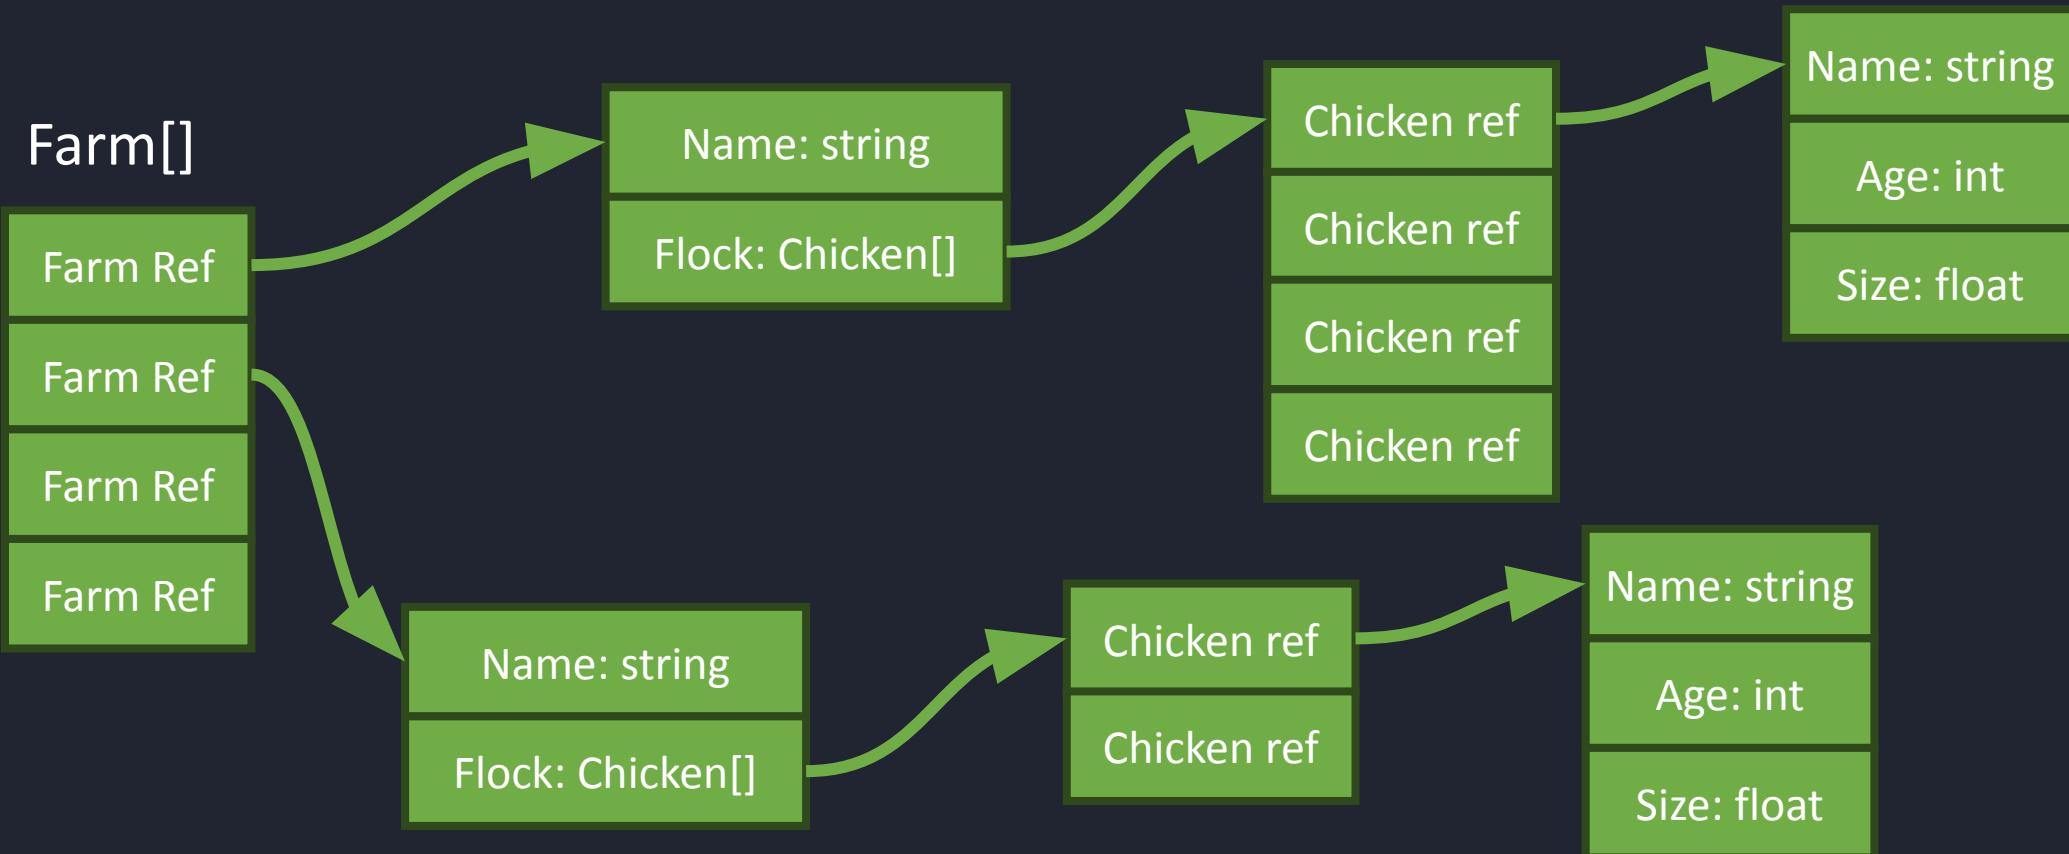
One to Many Relationships

```
type Chicken =
  {
    Name: string
    Size: float
    Age: int
  }

type Farm =
  {
    Name: string
    Flock: Chicken[]
  }
```

```
let cooperative = [ | // []<Farm>
  {
    Name = "Farm 1"
    Flock = [ |
      {
        Name = "Clucky"
        Age = 1
        Size = 1.0
      }
      {
        Name = "Drumstick"
        Age = 2
        Size = 1.0
      }
    ] ]
  }
```

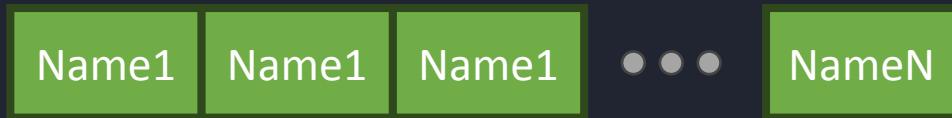
One to Many Relationships



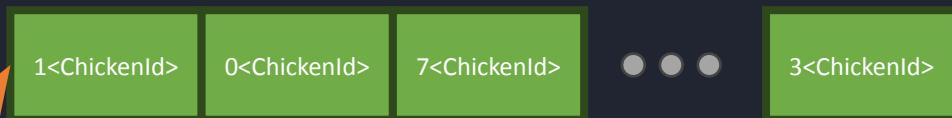
Farm Array Oriented Layout

Arrays for Farms

Names



Chickens

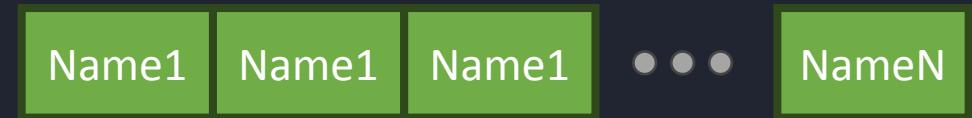


Ranges



Arrays for Chickens

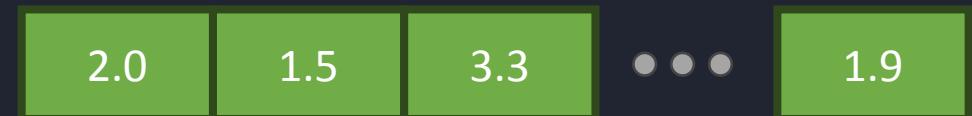
Names



Ages



Sizes



One to Many with Range and BarGroup

1<FarmlId>

Ranges

Range Start: 0 Length: 2	Range Start: 2 Length: 3	Range Start: 5 Length: 1	Range Start: 6 Length: 1
--------------------------------	--------------------------------	--------------------------------	--------------------------------

Chickens

1<ChickenId>	0<ChickenId>	3<ChickenId>	2<ChickenId>	5<ChickenId>	4<ChickenId>	6<ChickenId>	...
--------------	--------------	--------------	--------------	--------------	--------------	--------------	-----

Benchmark

- 100 Random Farms
- 100,000 Random Chickens randomly assigned to Farms
- Double the size of the Chickens

Benchmark

```
// * Summary *

BenchmarkDotNet v0.13.8, Windows 11 (10.0.22621.2134/22H2/2022Update/SunValley2)
13th Gen Intel Core i9-13900K, 1 CPU, 32 logical and 24 physical cores
.NET SDK 8.0.100-preview.7.23376.3
[Host]      : .NET 8.0.0 (8.0.23.37506), X64 RyuJIT AVX2 DEBUG
DefaultJob : .NET 8.0.0 (8.0.23.37506), X64 RyuJIT AVX2

| Method | Mean | Error | StdDev |
|-----:|-----:|-----:|-----:|
| ObjectStyle | 1,021.14 us | 22.755 us | 65.287 us |
| ArrayStyle | 24.01 us | 0.468 us | 0.481 us |
```

Benchmark

```
// * Summary *

BenchmarkDotNet v0.13.8, Windows 11 (10.0.22621.2134/22H2/2022Update/SunValley2)
13th Gen Intel Core i9-13900K, 1 CPU, 32 logical and 24 physical cores
.NET SDK 8.0.100-preview.7.23376.3
[Host]      : .NET 8.0.0 (8.0.23.37506), X64 RyuJIT AVX2 DEBUG
DefaultJob : .NET 8.0.0 (8.0.23.37506), X64 RyuJIT AVX2

| Method           | Mean       | Error    | StdDev   | Median     |
|-----:|-----:|-----:|-----:|-----:|
| ObjectStyle     | 980.407 us | 21.5794 us | 63.6272 us | 963.441 us |
| ArrayStyle      | 24.736 us  | 0.4845 us  | 0.5184 us  | 24.665 us  |
| ArraySIMDStyle  | 9.864 us   | 0.1066 us  | 0.0998 us  | 9.852 us  |
```

Range and BarGroup

```
[<Struct>]
type BarGroup<[<Measure>]> = {
    Measure: <Measure>;
    Value: <Value>;
    internal(ranges: Bar<'Measure, 'Range>, values: <'Value[]>) =
        // WARNING: This member is public for the purposes of inlining but it is not meant for public consumption.
        [<EditorBrowsable(EditorBrowsableState.Never)>]
        member __values = values // []<'Value>
        // WARNING: This member is public for the purposes of inlining but it is not meant for public consumption.
        [<EditorBrowsable(EditorBrowsableState.Never)>]
        member __ranges = ranges // Bar<'Measure, Range>

        member bg.Item // ReadOnlySpan<'Value>
            with inline get(i: int<'Measure>) =
                let ranges = bg.__ranges
                if i < ranges.Length then
                    let range = ranges[i]
                    ReadOnlySpan(bg.__values, range.Start, range.Length)
                else
                    ReadOnlySpan()
```

Other Tools

- BitSet
- ArrayPool
- StackAlloc
- ObjectPools
- StructPools

Memory Layout is Crucial for Performance

- The GC is your friend
- The GC doesn't understand your program's logic, just memory
- Ideal performance is only achieved by YOU laying out your data to be CPU-friendly

Predictably

Write logic that makes it easy for the CPU to predict logic and data access



Entity Component System (ECS)

- Entity
 - The elements of your program
 - Farms, Chickens, Users, etc.
- Component
 - The data associated with Entities
 - Location, Velocity, Status, etc.
- System
 - The logic that acts on Components
 - Movement, State Change, AI, etc.

Entities

```
module Ids =  
  
    []  
    type FarmId  
  
    []  
    type ChickenId  
  
    []  
    type TurkeyId  
  
    []  
    type GooseId  
  
    []  
    type AnimalId
```

Components

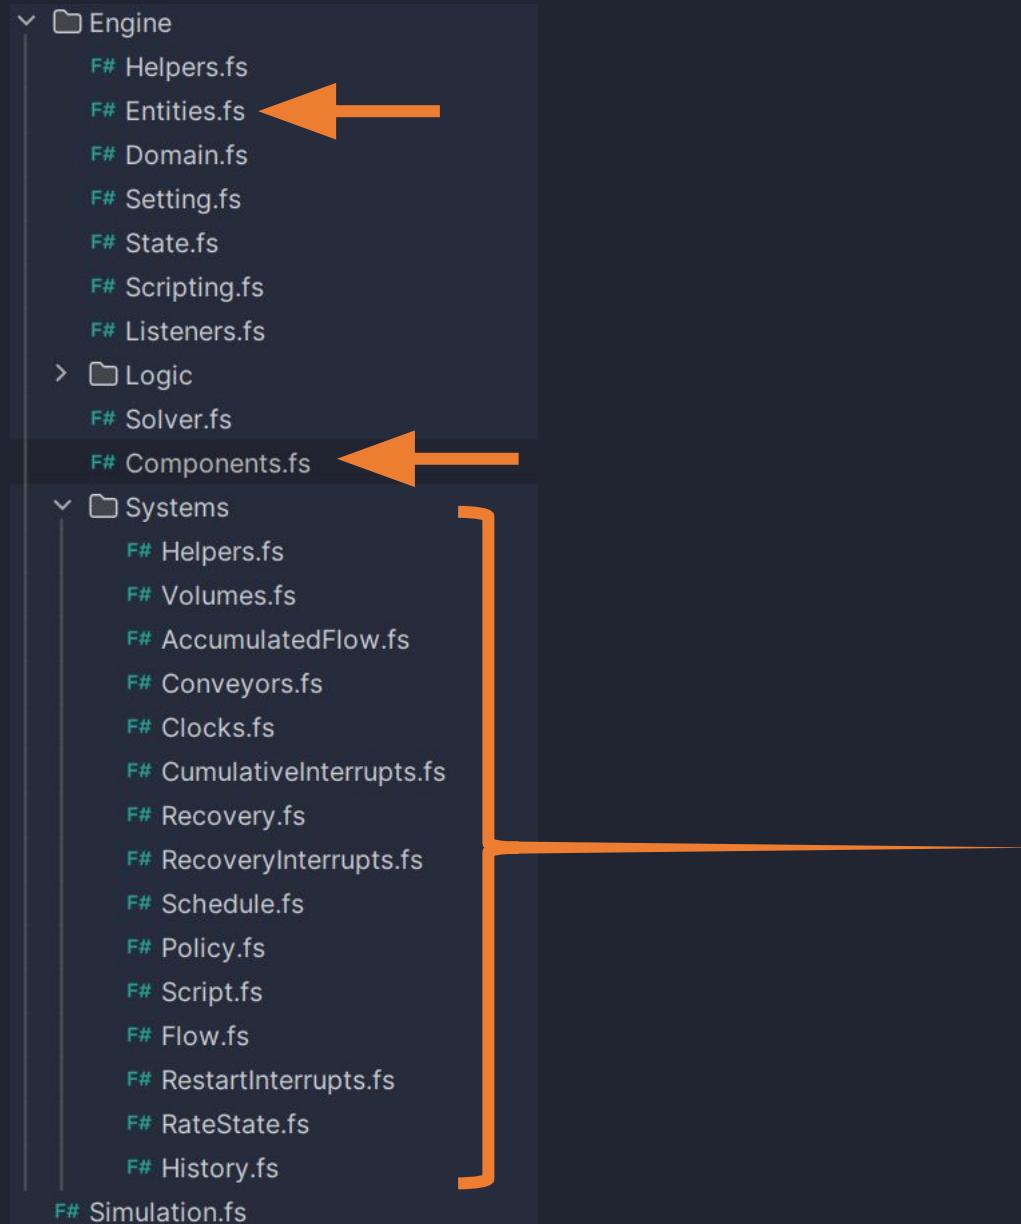
```
module Components =  
  
    type Flock =  
        {  
            Names: Bar<ChickenId, string>  
            Ages: Row<ChickenId, int>  
            Sizes: Row<ChickenId, float>  
        }  
  
        [  
            <NoComparison; NoEquality>  
        ]  
    type Rafter =  
        {  
            Names: Bar<TurkeyId, string>  
            Ages: Row<TurkeyId, int>  
            Sizes: Row<TurkeyId, float>  
            IsActive: BitSet<TurkeyId>  
        }  
  
    type Cooperative = {  
        Names: Bar<FarmId, string>  
        Chickens: BarGroup<FarmId, int<ChickenId>>  
        Turkeys: BarGroup<TurkeyId, int<TurkeyId>>  
    }
```

Systems

```
namespace Systems

module Chickens =
    let age // Row<ChickenId,int> -> Row<ChickenId,float> -> TimeSpan -> unit
        (ages: Row<ChickenId, int>)
        (sizes: Row<ChickenId, float>)
        (timeStep: TimeSpan)
        =
        // Logic for aging chickens
```

Systems



Why ECS for Predictability

- ECS Encourages you to group work by the data it acts on
- Systems act in bulk, which encourages small loops that are easy for CPUs to predict
 - Many Small Loops are faster than Fewer Large Loops



Real World Innovation

Proceedings of the 2008 Winter Simulation Conference

S. J. Mason, R. R. Hill, L. Mönch, O. Rose, T. Jefferson, J. W. Fowler eds.

DISCRETE RATE SIMULATION USING LINEAR PROGRAMMING

Cecile Damiron
Anthony Nastasi

6830 Via Del Oro, Suite 230
Imagine That Inc.
San Jose, CA 95119 USA

ABSTRACT

Discrete Rate Simulation (DRS) is a modeling methodology that uses event based logic to simulate linear continuous processes and hybrid systems. These systems are con-

ing its movement. Typically, these control mechanisms include valves for modulating flow rates and branching functions for merging and diverging the flow streams.

Hybrid systems, while similar to flow systems in that flow moves continuously at certain rates between loca-

Real World Innovation

Engine	Time for Simulation [sec]	Simulations per Hour
ExtendSim	2,685	1.34
Aidos	4	900

ExtendSim vs Aidos



Takeaways...

- F# lets you write correct code quickly
- F# can achieve C-levels of performance
- Units of Measure are essential
- The .NET Runtime is fast and getting faster

Thank you!

SDI

F# Community

Hack & Craft

Rheinland-Pfälzische Technische Universität
Kaiserslautern-Landau

FSLab

F# Software Foundation

Contact: hi@fastfsharp.com

